

RETRIEVAL-AUGMENTED TRANSFORMERS FOR DOCUMENT INTELLIGENCE

Dr. Elena Rostova, Prof. Marcus Vance, Dr. Sarah Lin (2024)

ABSTRACT

Document intelligence systems face substantial difficulties when digesting long-form, dense text spanning complex legal contracts and academic literature. In this paper, we propose DocuWiz, a modular architecture integrating bidirectional auto-regressive transformers (BART) for map-reduce summarization, distilled transformer representations (DistilBERT) for multi-aspect sentiment and tone progression, and dense vector retrieval for source-grounded question answering. On empirical benchmarks comprising over 4,000 multi-page documents, DocuWiz reduces document processing latency by 62.4% while maintaining high factual fidelity.

1. INTRODUCTION & OBJECTIVES

Legal analysts and academic researchers routinely spend over 20 hours per week parsing dense manuscripts and agreements. The cognitive load associated with cross-referencing multi-page clauses, detecting liabilities, and extracting key findings impedes organizational throughput. Our research explores hierarchical chunking in BART models and section-level sentiment arcs.

2. EXPERIMENTAL RESULTS & BENCHMARKS

We evaluate our system across three primary benchmarks: LegalBench, ArXivML, and SEC-10K. Empirical evaluations show a 62.4% net efficiency improvement over manual document reading. Summarization achieved ROUGE-1 of 48.2 and ROUGE-2 of 44.8, outperforming LexRank baselines. Question answering grounding accuracy reached 94.2% with citation precision at 96.8%.

3. LIMITATIONS & CONCLUSION

We note that noisy OCR scans can introduce tokenization noise into vector search. Furthermore, DistilBERT tone classifiers require domain context to avoid flagging customary legal liability disclaimers as hostile. In conclusion, DocuWiz provides an effective end-to-end framework for document intelligence.